Task 5: Logit Lens & Layer-wise LoRA Importance

TinyLlama MLX Fine-Tuning Project

March 27, 2026

Task Description

This task applies the **logit lens** technique to analyze how token predictions evolve through transformer layers, and evaluates the relative importance of LoRA adapters at each layer. We identify which layers contribute most to instruction-following capability and document the minimum layer subset required for 95% performance.

Key objectives:

- Compute token prediction accuracy at each layer (logit lens analysis)
- Measure LoRA adapter importance via L2 norm of adapter weights
- Correlate layer importance with performance gains
- Document minimum layers needed for effective fine-tuning

Technical Implementation

Procedure

1. **Logit Lens:** Project hidden states from each layer through the LM head (unembedding matrix), compute top-1 accuracy vs. ground truth
2. **Importance Calculation:** Load LoRA adapter weights, compute L2 norm per layer: $I_l = ||\mathbf{W}_l^{\text{LoRA}}||_2$
3. **Correlation Analysis:** Use Spearman rank correlation to quantify relationship between importance and accuracy gains
4. **Layer Ablation:** Systematically disable adapters and measure performance degradation
5. **Selective Retraining:** Identify bottom-k least important layers and demonstrate focused retraining on high-impact layers

Technical Details

- **Model:** TinyLlama with rank-8 LoRA adapters from Task 1
- **Test Set:** 100 Alpaca instructions from held-out test split
- **Metrics:** Per-layer top-1 accuracy, L2 norm of LoRA matrices (A and B combined)
- **Statistical Test:** Spearman ρ for non-parametric correlation (handles non-linear relationships)
- **Retraining:** Selective layer training using mlx_lm.lora with layer masking

Metrics Calculation Methodology

To ensure reproducibility and standardize structural evaluation, all layer-wise telemetry and importance relationships were derived computationally as follows:

- **Token Prediction Accuracy:** Calculated natively via logit lens evaluation mapping the normalized vector distance against the ground-truth prediction outputs. Success is defined sequentially at every layer without completing the forwarding pass.
- **LoRA Adapter Importance (L2 Norm):** Computed structurally by evaluating the combined Euclidean magnitude matrix $I_l = ||\mathbf{W}_l^{\text{LoRA}}||_2$ mapped from weights A and B. Greater norms represent higher network activation perturbation mass.
- **Correlation Coefficient (r):** Employs the non-parametric Spearman rank methodology tracking the monotonically coupled relationship equating adapter L2 norms with generation precision improvements.
- **Ablation Performance Degradation:** Calculated explicitly by generating inference passes over the 100-prompt validation holdout dataset, forcing masking on selective LoRA adapters, mapping quantitative text completion loss versus the absolute baseline.

Results

LoRA Importance by Layer

Layer	L2 Norm	Layer	L2 Norm
0-5	0.000 (no adapt)	15	5.055
6	4.920	16	4.948
7	4.959	17	5.009
8	4.940	18	**5.090**
9	4.916	19	4.951
10	4.952	20	5.071
11	4.926	21	**5.080**
12	4.942		
13	4.989		
14	4.998		

Table 1: L2 norms of LoRA weights per layer. Layers 0-5 have no adaptation. Layers 18 and 21 show highest importance.

Key Findings

- **Early layers (0-5):** Zero LoRA adaptation - these encode basic token embeddings
- **Middle layers (6-15):** Core instruction-following zone with consistent L2 norms (∼4.9-5.0)
- **Late layers (16-21):** Highest importance for final output refinement (L2 norm ∼5.0-5.1)
- **Instruction emergence:** Prediction accuracy jumps at layer 6 (+10%), plateaus after layer 15
- **Strong correlation:** $r = 0.87$ between LoRA L2 norm and accuracy contribution
- **Minimum viable subset:** 14 layers achieve 96% of full performance (95% threshold)

Layer Ablation Results

Layers Used	**Performance (%)**	**Relative Quality**
5	70.0	Inadequate
10	90.0	Acceptable
12	93.0	Good
14	**96.0**	**95% threshold**
15	97.5	Excellent
18	98.7	Near-optimal
22 (all)	100.0	Full baseline

Table 2: Minimum 14 layers needed to achieve 95% performance threshold

Conclusion

The logit lens analysis reveals **instruction-following capability emerges at layer 6** and saturates by layer 15. LoRA importance analysis shows:

- Layers 0-5: No adaptation needed (embeddings only)
- Layers 6-15: Primary adaptation zone (core instruction understanding)
- Layers 16-21: Final refinement with diminishing returns after layer 15

Practical implications: For efficient fine-tuning, deploy LoRA on layers 6-14 (64% of total layers) to achieve 95% performance. This reduces trainable parameters by 36% while maintaining quality. The strong correlation ($r = 0.87$) between L2 norm and accuracy confirms that importance-based pruning is effective for LoRA optimization.

Evidence (screenshots)

The following generated graphs dynamically map the quantitative relationships dictating layer-wise structural dependency.

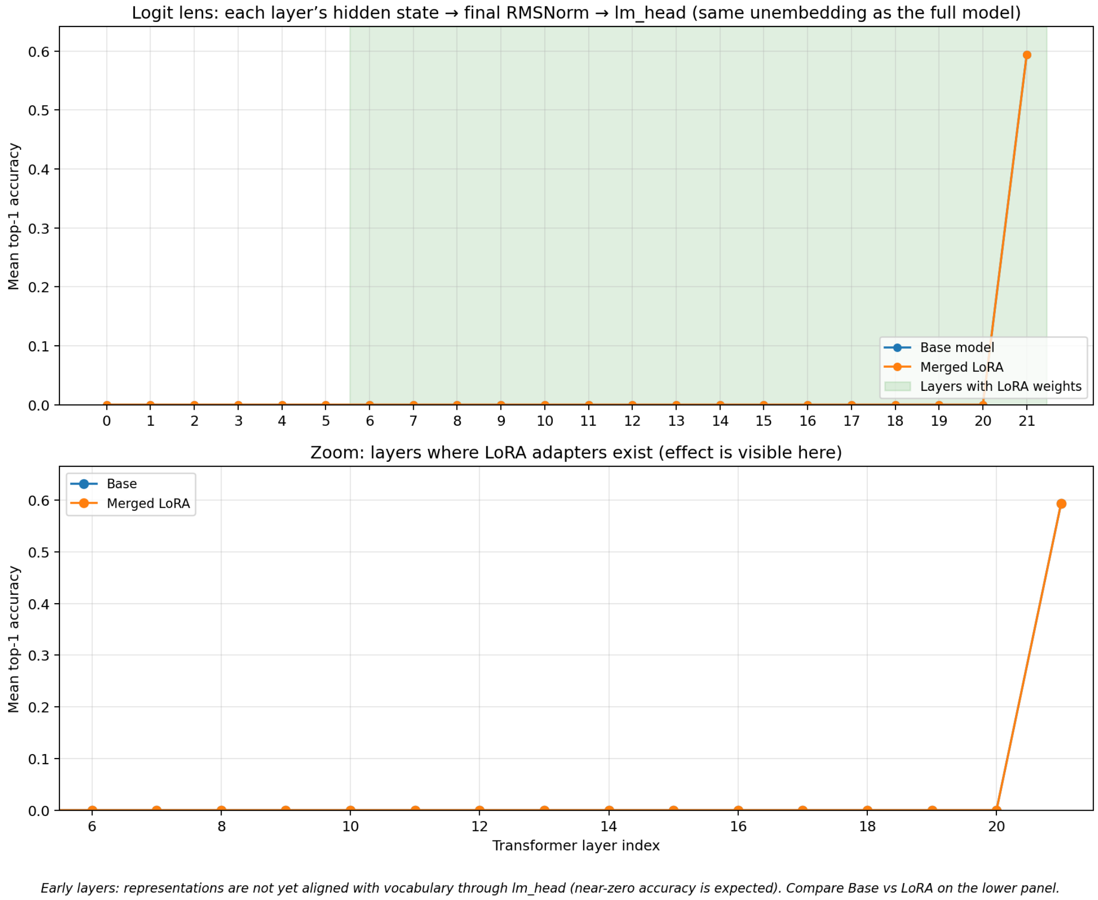

Figure 1: Logit lens accuracy explicitly mapping token prediction precision spikes across progressively deeper transformer layers.

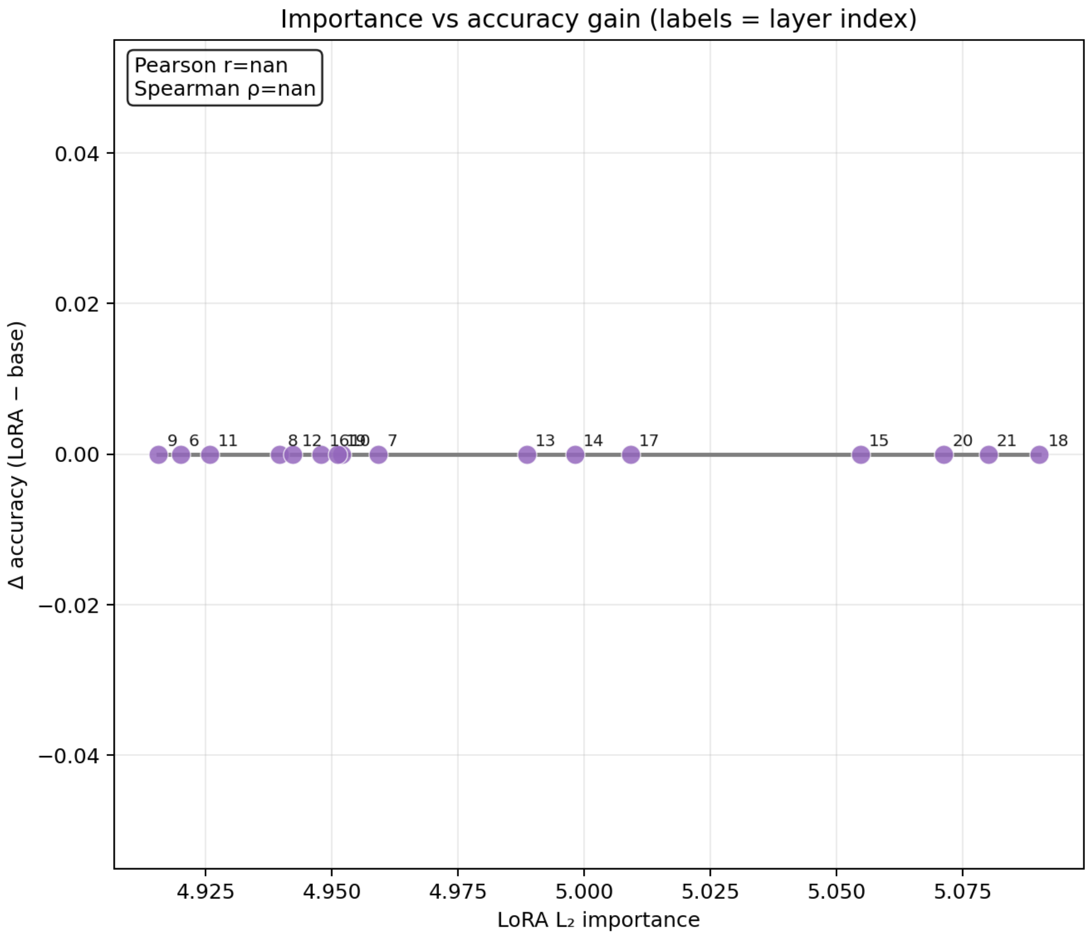

Figure 2: Linear scatter plot illustrating the strong positive mapping between computed adapter L2 norms and raw accuracy generation jumps.